

```
!pip install keras
import keras
```

```
import tensorflow as tf
tf.test.gpu_device_name()
```

```
!unzip -o -q BoatOrCar.zip
```

In [4]:

```
import numpy as np
from keras.models import Sequential
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense
```

In [6]:

```
from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale = 1./255)
training_set = train_datagen.flow_from_directory('DataBoatOrCar/training',
                                                target_size = (128, 128),
                                                batch_size = 32,
                                                class_mode = 'binary')

test_datagen = ImageDataGenerator(rescale = 1./255)
test_set = test_datagen.flow_from_directory('DataBoatOrCar/test',
                                            target_size = (128, 128),
                                            batch_size = 32,
                                            class_mode = 'binary')

from matplotlib import pyplot as plt
x, y = training_set.next()
for i in range(0, 10):
    image = x[i]
    plt.imshow(image)
    plt.show()
```

Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

<Figure size 640x480 with 1 Axes>

In [9]:

```

classifier = Sequential()

classifier.add(Conv2D(32, (3, 3), input_shape = (128, 128, 3), activation = 'relu'))

classifier.add(MaxPooling2D(pool_size = (2, 2)))

classifier.add(Conv2D(32, (3, 3), activation = 'relu'))

classifier.add(MaxPooling2D(pool_size = (2, 2)))

classifier.add(Conv2D(32, (3, 3), activation = 'relu'))

classifier.add(Flatten())

classifier.add(Dense(units = 128, activation = 'relu'))
classifier.add(Dense(units = 128, activation = 'relu'))
classifier.add(Dense(units = 128, activation = 'relu'))

classifier.add(Dense(units = 1, activation = 'sigmoid'))

classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])

classifier.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_5 (Conv2D)	(None, 61, 61, 32)	9248
max_pooling2d_4 (MaxPooling2D)	(None, 30, 30, 32)	0
conv2d_6 (Conv2D)	(None, 28, 28, 32)	9248
flatten_2 (Flatten)	(None, 25088)	0
dense_5 (Dense)	(None, 128)	3211392
dense_6 (Dense)	(None, 128)	16512
dense_7 (Dense)	(None, 128)	16512
dense_8 (Dense)	(None, 1)	129
Total params: 3,263,937		
Trainable params: 3,263,937		
Non-trainable params: 0		

In [10]:

```

history = classifier.fit_generator(training_set,
                                  steps_per_epoch = 2000,
                                  epochs = 5,
                                  validation_data = test_set,
                                  validation_steps = 1000)

classifier.save("BoatOrCarModel.hdf5") # Hierarchical Data Format Fifth Version
print("Saved model to disk")

import pickle
filehandler = open('history.pkl', 'wb')
pickle.dump(history, filehandler)

```

Epoch 1/5

2000/2000 [=====] - 1351s 676ms/step - loss: 0.0642 - accuracy: 0.9728 - val_loss: 1.1397 - val_accuracy: 0.8551

Epoch 2/5

2000/2000 [=====] - 1348s 674ms/step - loss: 7.4446e-04 - accuracy: 0.9993 - val_loss: 2.1193 - val_accuracy: 0.8557

Epoch 3/5

2000/2000 [=====] - 1358s 679ms/step - loss: 0.0181 - accuracy: 0.9936 - val_loss: 0.3384 - val_accuracy: 0.8173

Epoch 4/5

2000/2000 [=====] - 1343s 671ms/step - loss: 0.0038 - accuracy: 0.9986 - val_loss: 0.1763 - val_accuracy: 0.8210

Epoch 5/5

2000/2000 [=====] - 1345s 673ms/step - loss: 7.1588e-04 - accuracy: 0.9993 - val_loss: 3.0874 - val_accuracy: 0.8171

Saved model to disk

```
from keras.models import load_model
classifier = load_model('BoatOrCarModel.hdf5')

print(classifier.evaluate(training_set))

print(classifier.evaluate(test_set))

test_set.reset()
y_pred = classifier.predict_generator(test_set)
y_test = test_set.classes[test_set.index_array]
y_pred = ((y_pred>0.5)*1.0).flatten()
print(sum(y_pred==y_test)/y_pred.size)

!pip install sklearn
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

$$\begin{bmatrix} 85 & 415 \end{bmatrix}$$

In [12]:

```
import pickle
with open('history.pkl', 'rb') as pickle_file:
    content = pickle.load(pickle_file)

import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']



epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```



In []: